

# Security in and for Operating Systems

Jun Li

[lijun@cs.uoregon.edu](mailto:lijun@cs.uoregon.edu)

# Learning Objectives

- Security methods of ordinary OS
  - Memory and address protection
  - Control of access to general objects
  - File protection
  - User authentication
- Trusted OS
  - Concepts
  - Security policies and models
  - Trusted OS design:
    - kernelized design, separation/isolation, virtualization

# Security Methods of Ordinary OS

# Protected Objects

- Memory
- I/O devices (sharable, serially usable)
- Networks
- Sharable programs and subprocedures
- Sharable data

# Security Methods of OS

- Separation: One user's objects separate from other users
- Physical separation
  - E.g., separate printers
- Temporal separation
  - Processes executing at different times
- Logical separation
  - Users operate under the illusion that no other processes exist
- Cryptographic separation
  - Data and computations unintelligible to outside processes

# Levels of Separation

more difficult  
↓

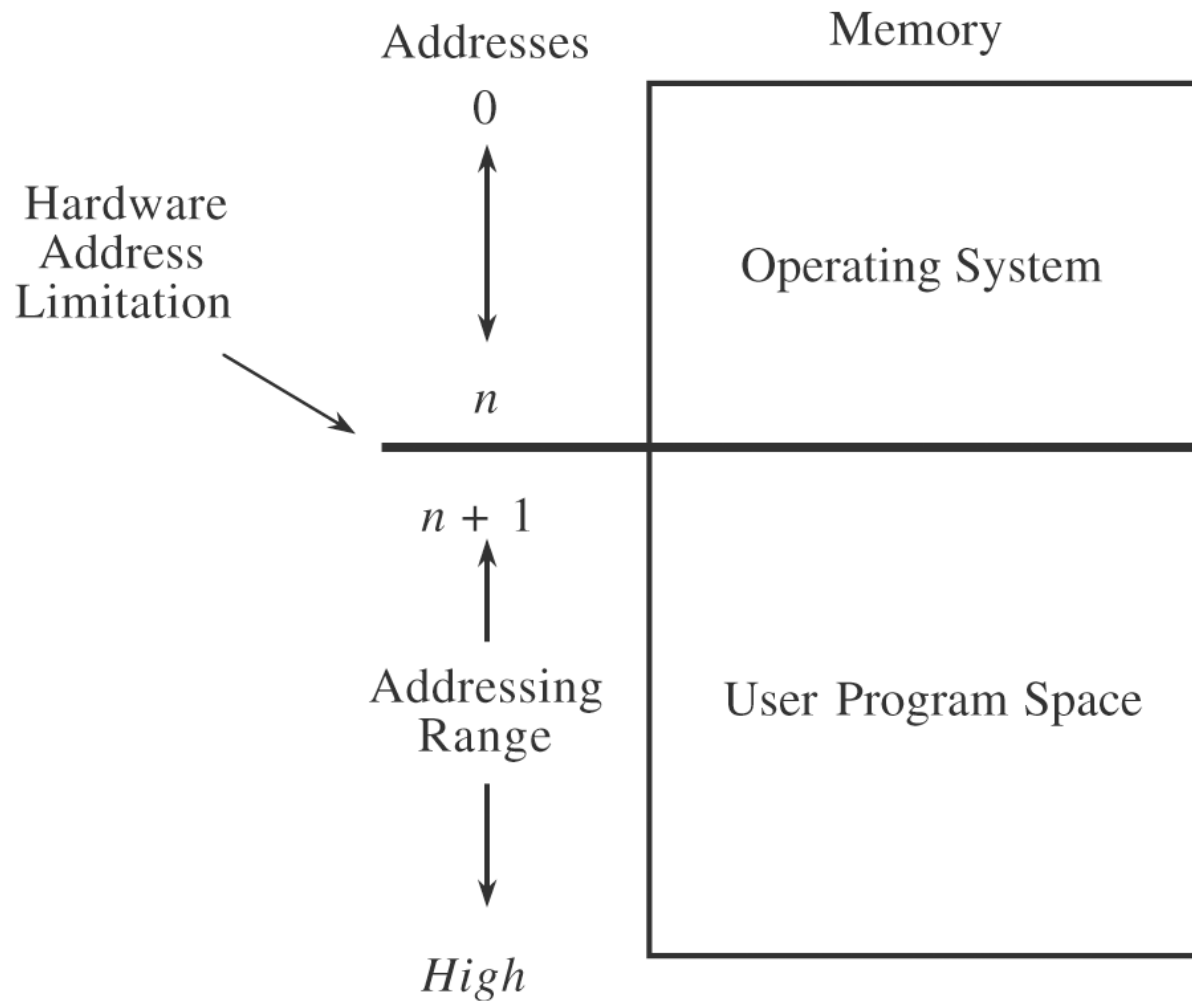
- Do not protect
- Isolate: Processes are unaware of each other
  - Each has its own address space, files, and other objects
- Share all or nothing
  - The owner of an object make it either public or private
- Discretionary access control
  - The owner controls the access to its objects
- Mandatory access control
  - The O.S. controls the access to objects
- Limit use of an object
  - Not just the access, but also the usage made after access

# Memory and Address Protection

Every access goes through certain HW points

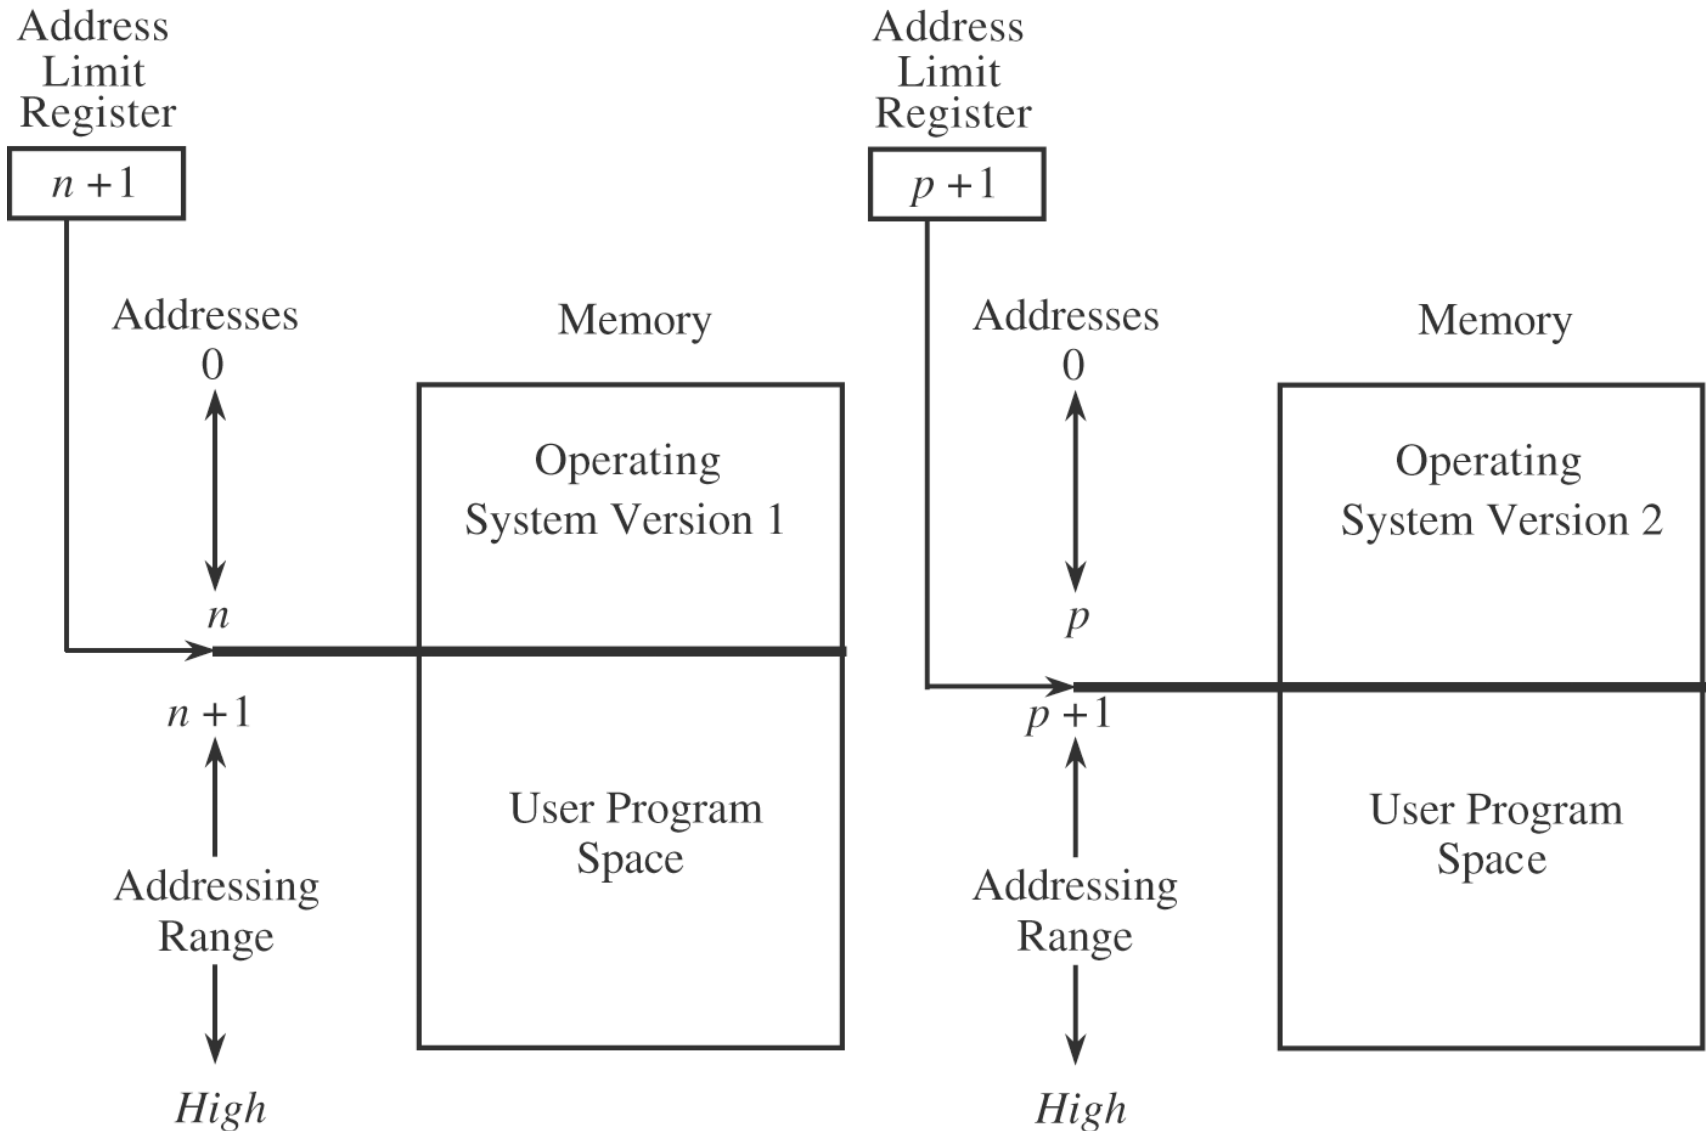
- Fence
- Relocation
- Base/Bound Registers
- Tagged Architecture
- Segmentation
- Paging
- Combined Paging with Segmentation

# Fence - Predefined





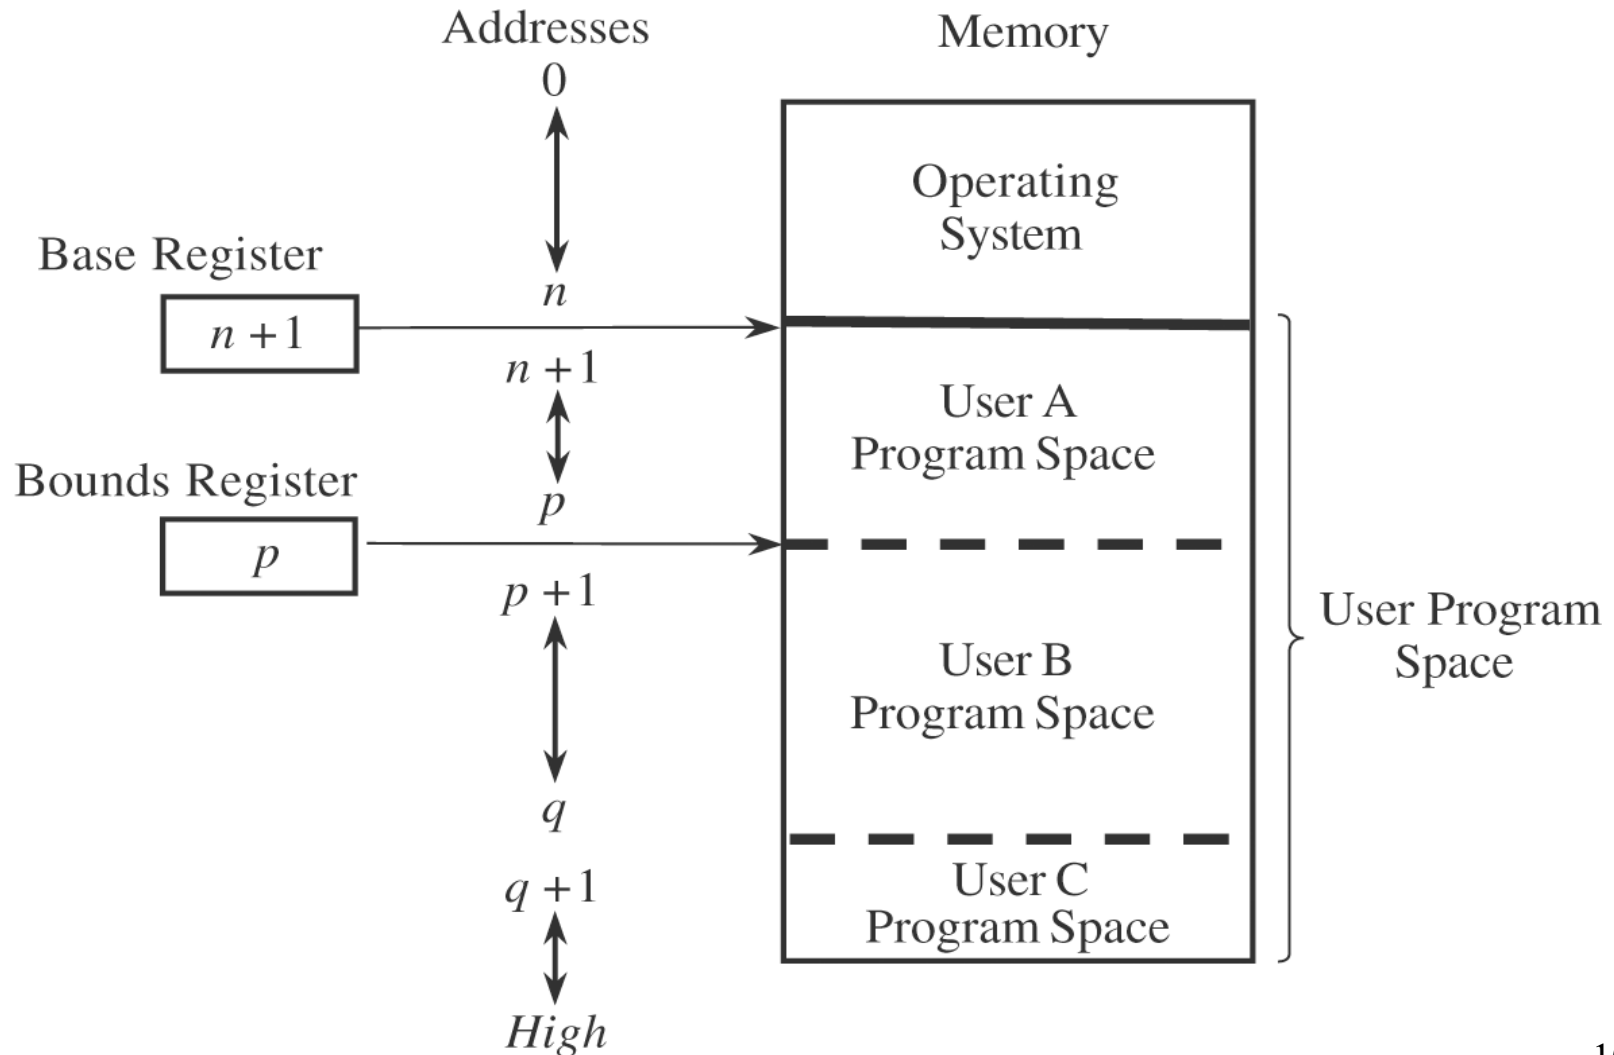
# Fence – Fence Register

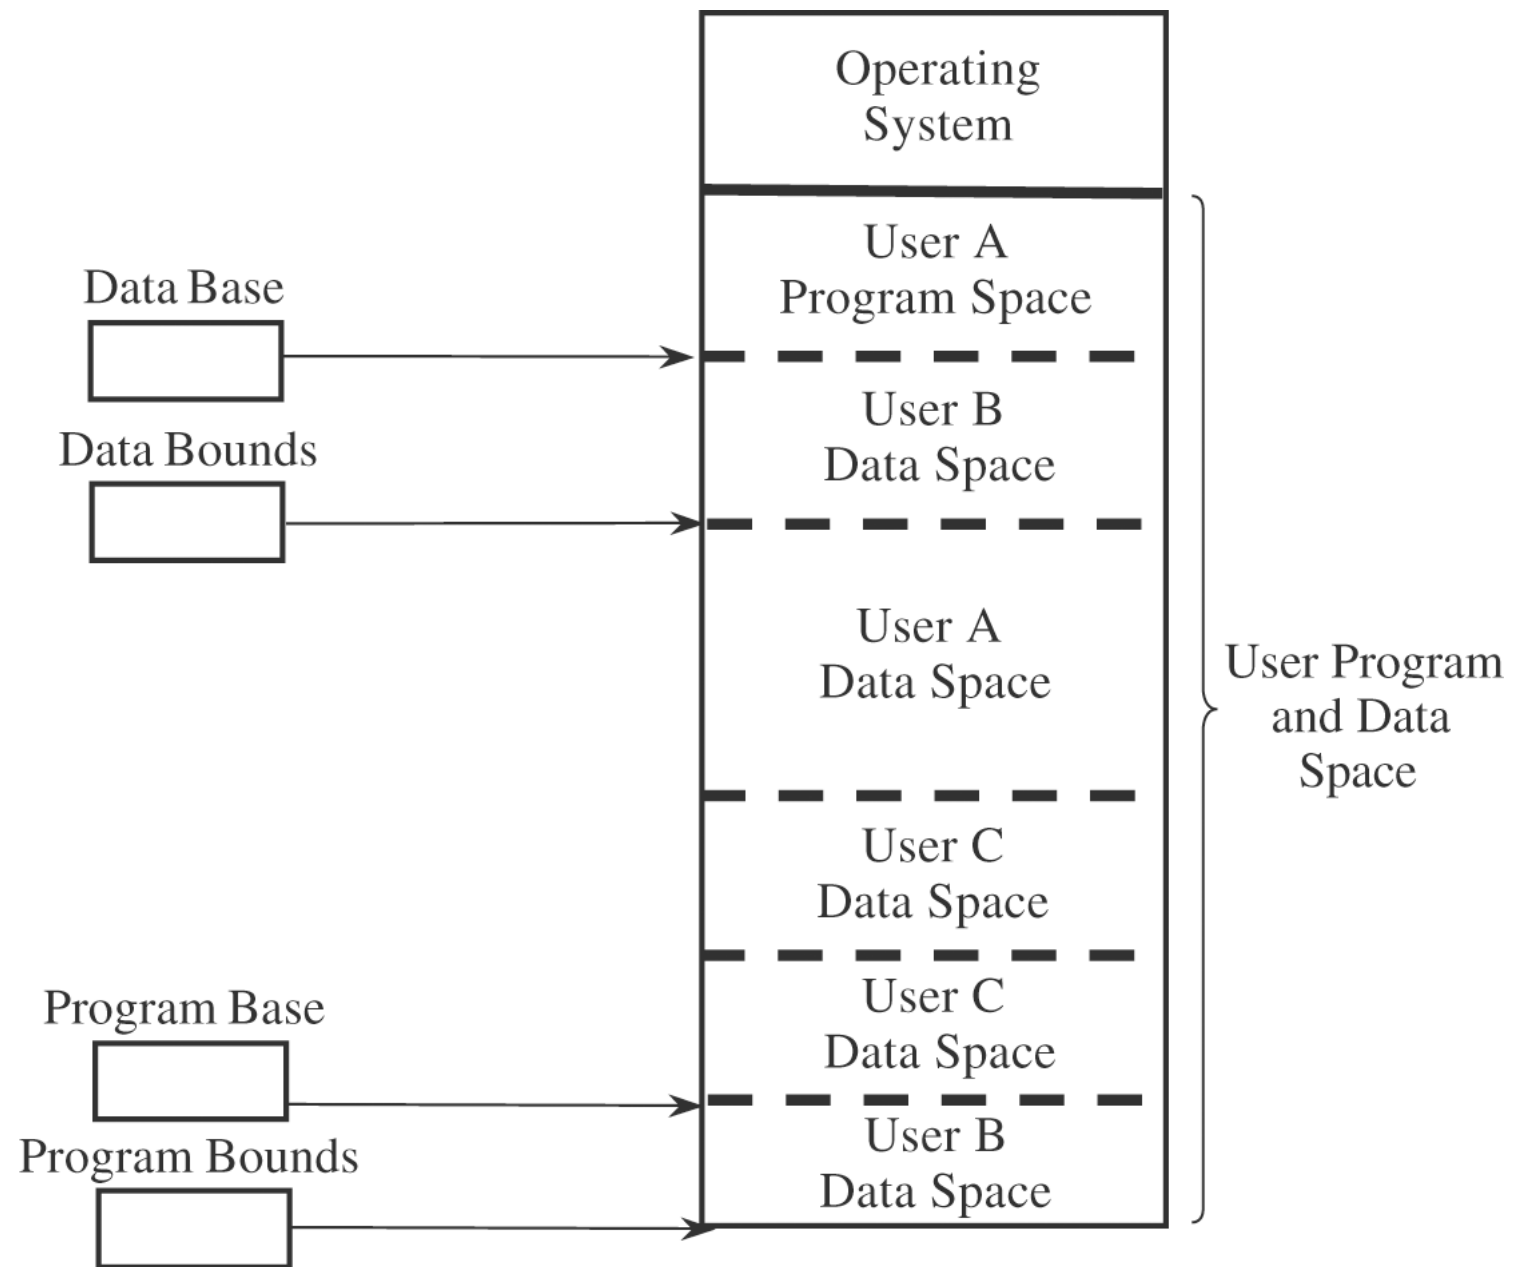


# Relocation







- Program written as if it began at address 0
- Adding a constant relocation factor when loading the program into memory
- The fence register can be a hardware relocation device

# Base/Bounds Registers



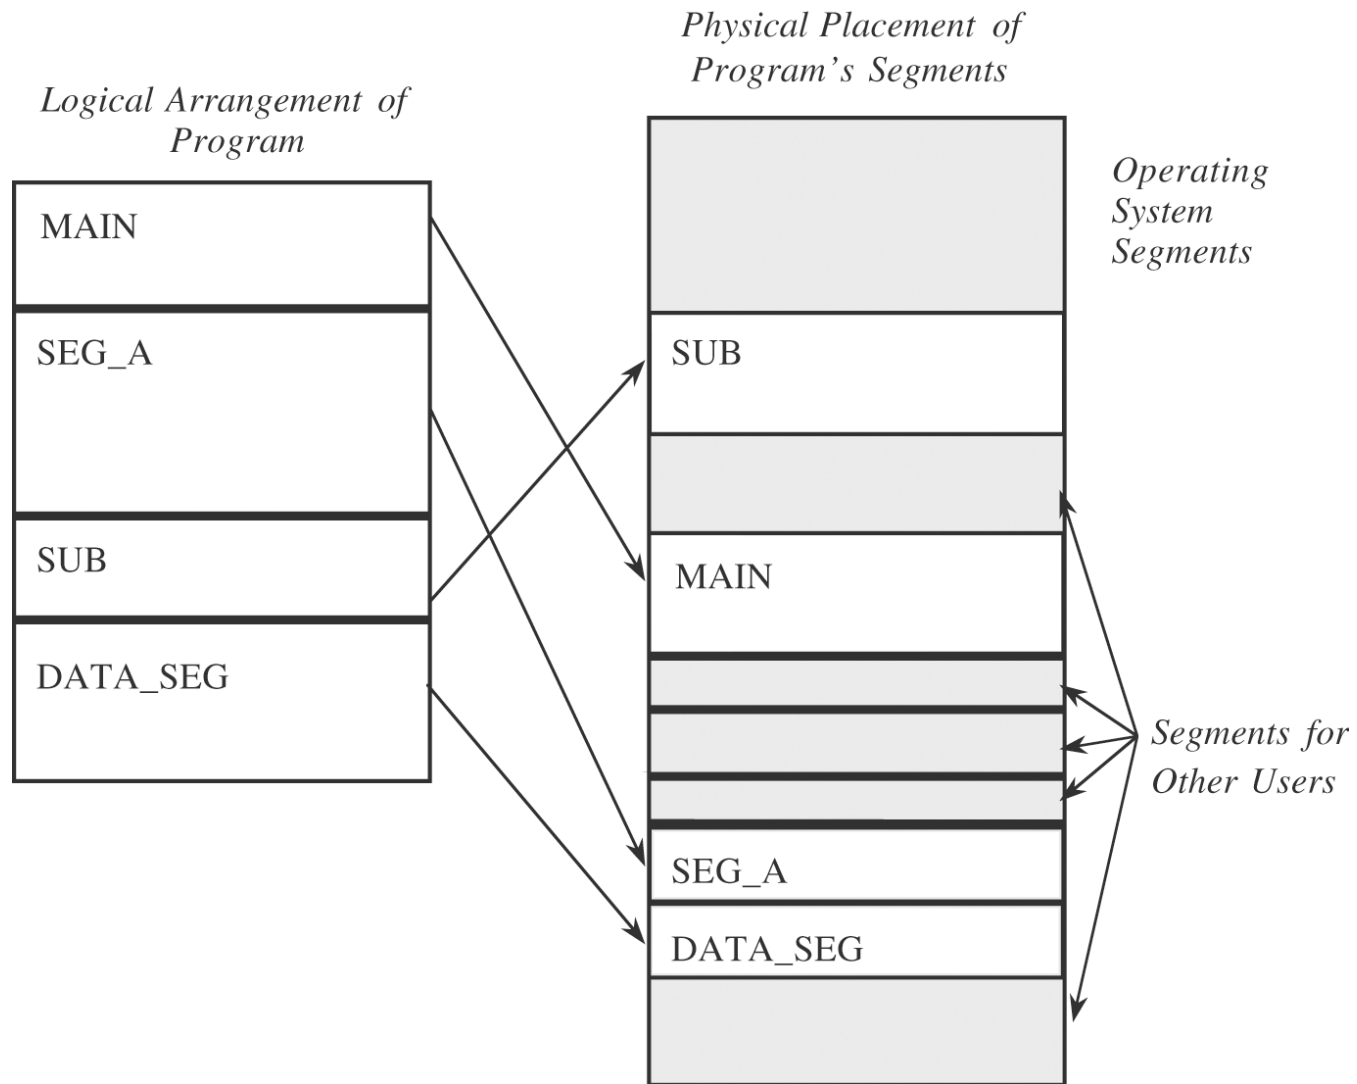


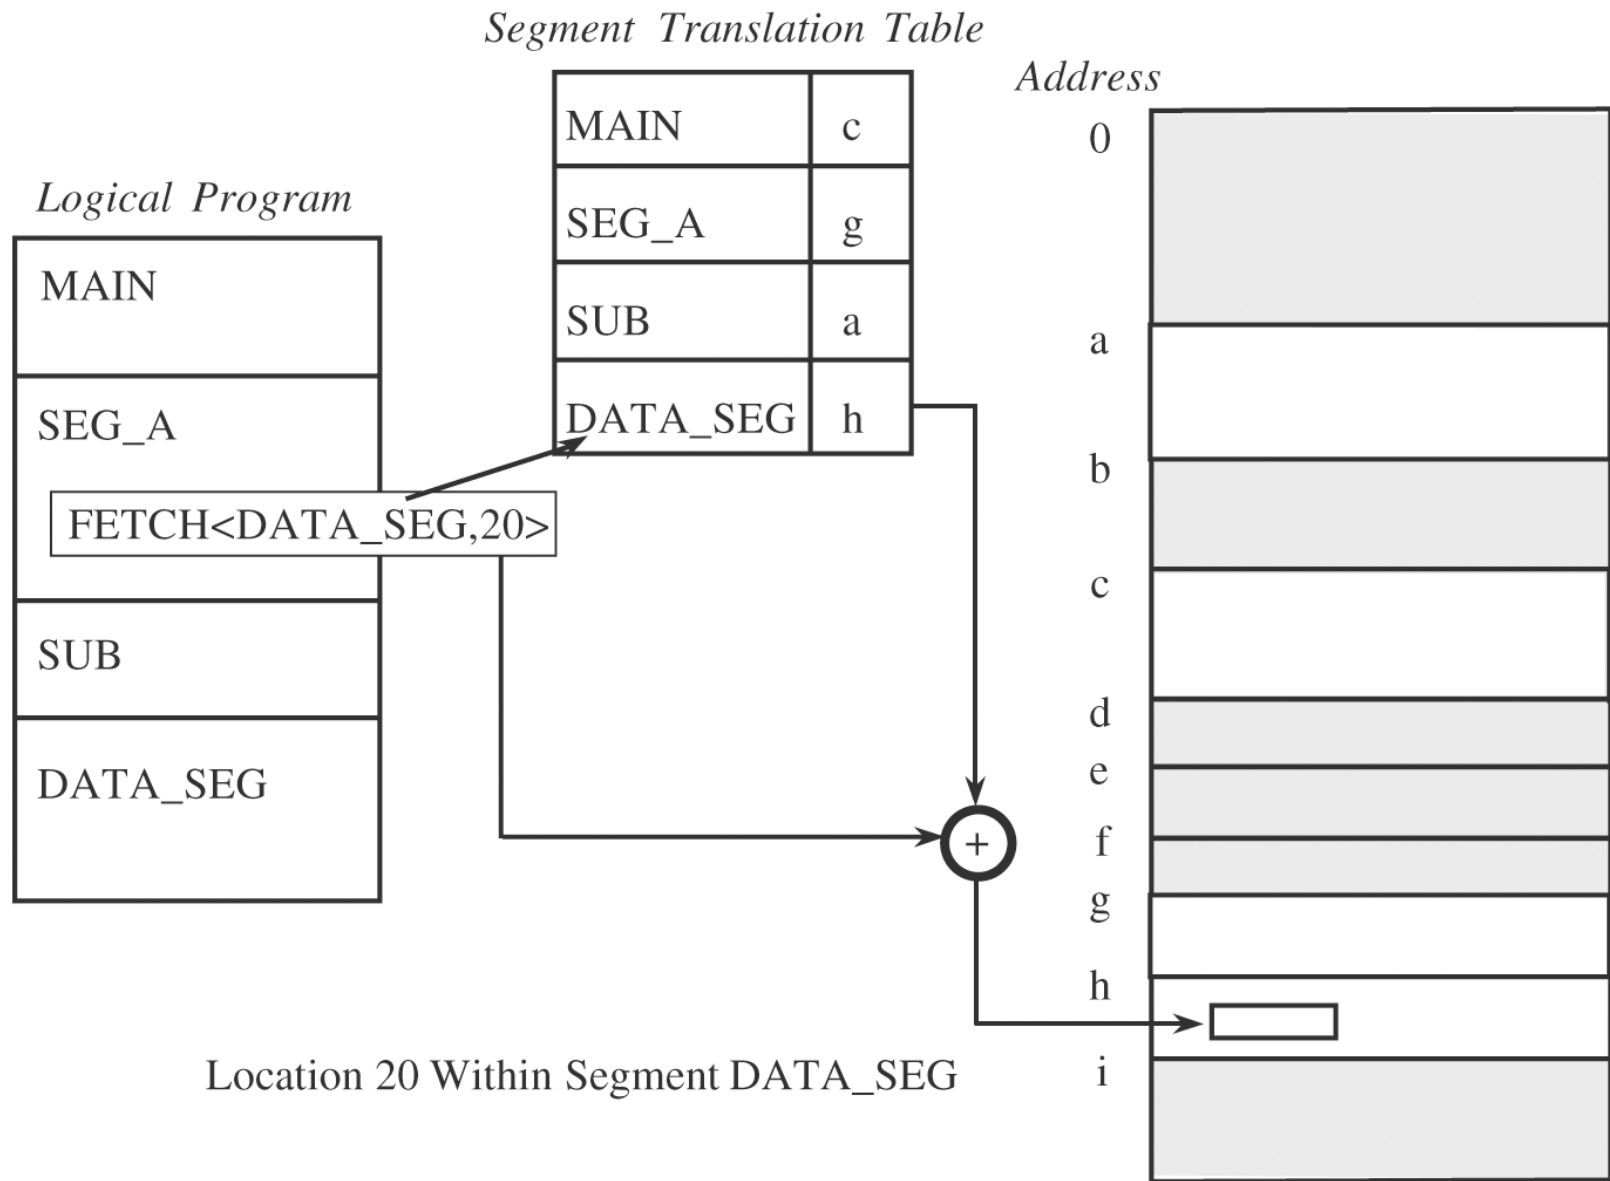
# Tagged Architecture

Tag	Memory Word
R	0001
RW	0137
R	0099
X	
X	
X	
X	
X	
X	
R	4091
RW	0002

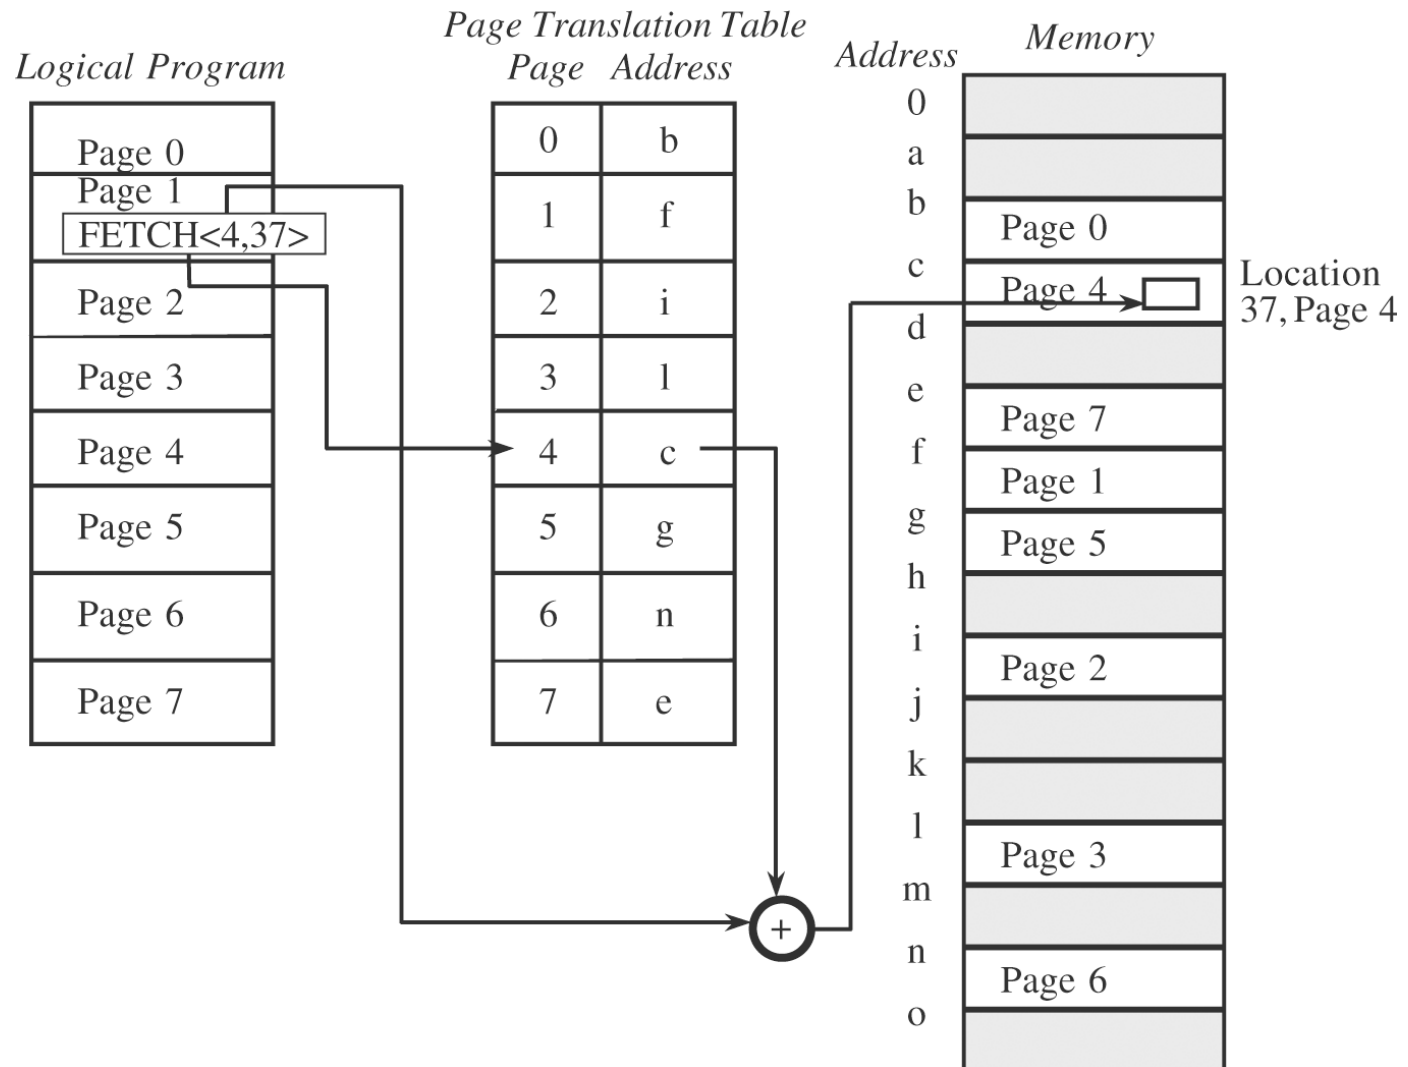
Code: R = Read-only    RW = Read/Write  
X = Execute-only

# Segmentation



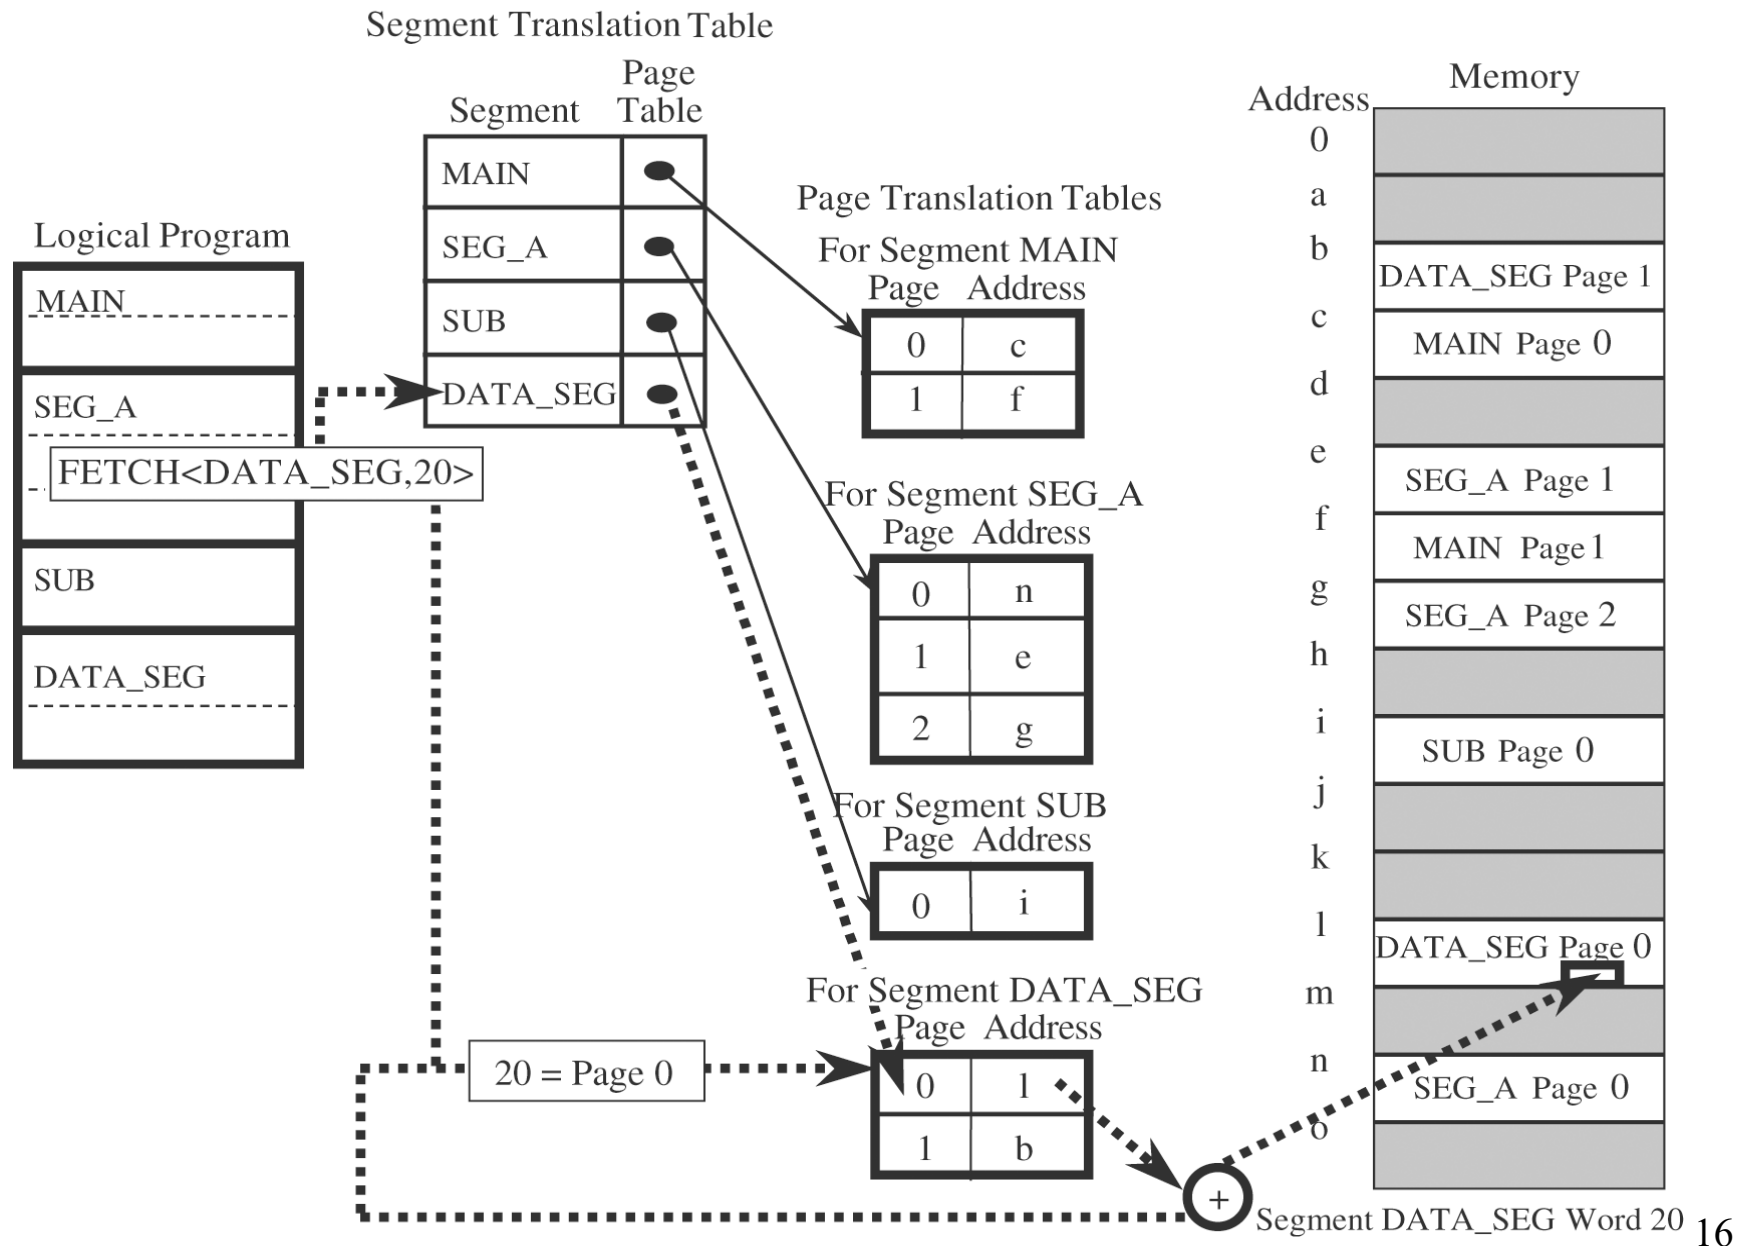


# Paging





# Combined Paging with Segmentation



# Control of Access to General Objects

- Memory is a special object
- Objects can be:
  - A memory space
  - A file or data set on an auxiliary storage device
  - An executing program in memory
  - A directory of files
  - A hardware device
  - A data structure, such as a stack
  - A table of the operating system
  - Privileged instructions
  - Passwords and user authentication mechanism
  - The protection mechanism itself
- Refer to Chapter 2 on Authentication and Access Control

# Designing Trusted Operating Systems

# 4 Underpinnings of a *Trusted* OS

- **Security policy:** requirement statements of what an OS should do and how it should do it, usually a set of rules on what to be secured and why.
- **Model:** a representation of the policy the OS will enforce (and compared with the security policy to ensure the security needs are met)
- **Design:** design/implement the OS w.r.t the model
- **Trust:** *features* (the OS has all the functionalities to enforce the security policy) and *assurance* (it will do so correctly and effectively)

# Secure vs. Trusted

- *Either-or*: something either is or is not secure
- Property of *presenter*
- *Asserted* based on product characteristics
- *Absolute*: not qualified as to how used, where, when, or by whom
- *A goal*
- *Graded*: degrees of trustworthiness
- Property of *receiver*
- *Judged* based on evidence and analysis
- *Relative*: viewed in context of use
- *A characteristic*

# Security Policies

- Military Security Policy
- Commercial Security Policies
  - Clark-Wilson
  - Separation of Duty
  - Chinese Wall

# Military Security Policy: object

- Each piece of information is ranked at a sensitivity level: *unclassified, restricted, confidential, secret, top secret*
- Each piece of information is also associated with *compartments*
  - Need-to-know principle
- **class or classification** of a piece of information: *<rank; compartments>*

# Military Security Policy: subject

- **clearance** of a subject:  $\langle rank; compartments \rangle$ 
  - The subject is trusted to access information up to a certain level of sensitivity (*rank*)
  - The subject needs to know certain categories of information (compartments)
- A subject  $s$  dominates an object  $o$  ( $s \geq o$ ) *iff*  
 $rank(s) \geq rank(o)$  *and*  
 $Compartments(s) \geq Compartments(o)$



# Military Security Policy

- A subject  $s$  can read an object  $o$  only if  $s$  dominates  $o$  ( $s \geq o$ )
- Appropriate for a setting in which access is rigidly controlled by a central authority

# Models of Security

- Test a particular policy
  - Completeness
  - Consistency
- Document a policy
- Help conceptualize and design an implementation
- Check whether an implementation meets its requirements

# Models

- Models for multilevel security
  - Bell-La Padula Confidentiality Model
  - Biba Integrity Model
- Models proving theoretical limitations of security systems
  - Graham-Denning
  - Harrison-Ruzzo-Ullman
  - Take-Grant
- We focus on models for multilevel security

# Bell-La Padula Confidentiality Model

- The system has a set of subjects  $S=\{s\}$  and a set of objects  $O=\{o\}$
- Each subject  $s$  or object  $o$  has a **security class**  $C(s)$ ,  $C(o)$
- Enforces two properties on the secure flow of information

# Properties

- **Simple Security Property.** A subject  $s$  can read an object  $o$  only if  $C(s) \geq C(o)$ .
  - Prevents *read-up*
- **\*-property** (called “star property”). A subject  $s$  who can read an object  $o$  can write to  $p$  only if  $C(o) \leq C(p)$ .
  - Prevents *write-down*

# Biba Integrity Model

- To prevent inappropriate modification of data
- Subjects and objects are ordered by an integrity classification scheme.  $I(s)$ ,  $I(o)$ .
- **Simple integrity property.** Subject  $s$  can modify (write) object  $o$  only if  $I(s) \geq I(o)$ .
- **Integrity \*-property.** If subject  $s$  has *read* access to object  $o$  with integrity level  $I(o)$ ,  $s$  can write object  $p$  only if  $I(o) \geq I(p)$ .

# Graham-Denning Model

- Model the access control mechanisms of a protection system
- A set of subjects  $S$ , a set of objects  $O$ , a set of rights  $R$ , and an access control matrix  $A$ .
- 8 primitive protection rights
  - Create object, create subject, delete object, delete subject
  - Read access right, grant access right, delete access right, transfer access right

# Harrison-Ruzzo-Ullman

- A variation of the Graham-Denning model
- Based on **commands** with conditions and primitive operations

```
command name( $o_1, o_2, \dots, o_k$ )  
    if       $r_1$  in  $A[s_1, o_1]$  and  
            $\dots$   
            $r_m$  in  $A[s_m, o_m]$   
    then  
            $op_1$   
            $\dots$   
            $op_n$   
end
```

- Primitive operations:
  - create/destroy subject/object
  - enter/delete right  $r$  into/from  $A[s, o]$

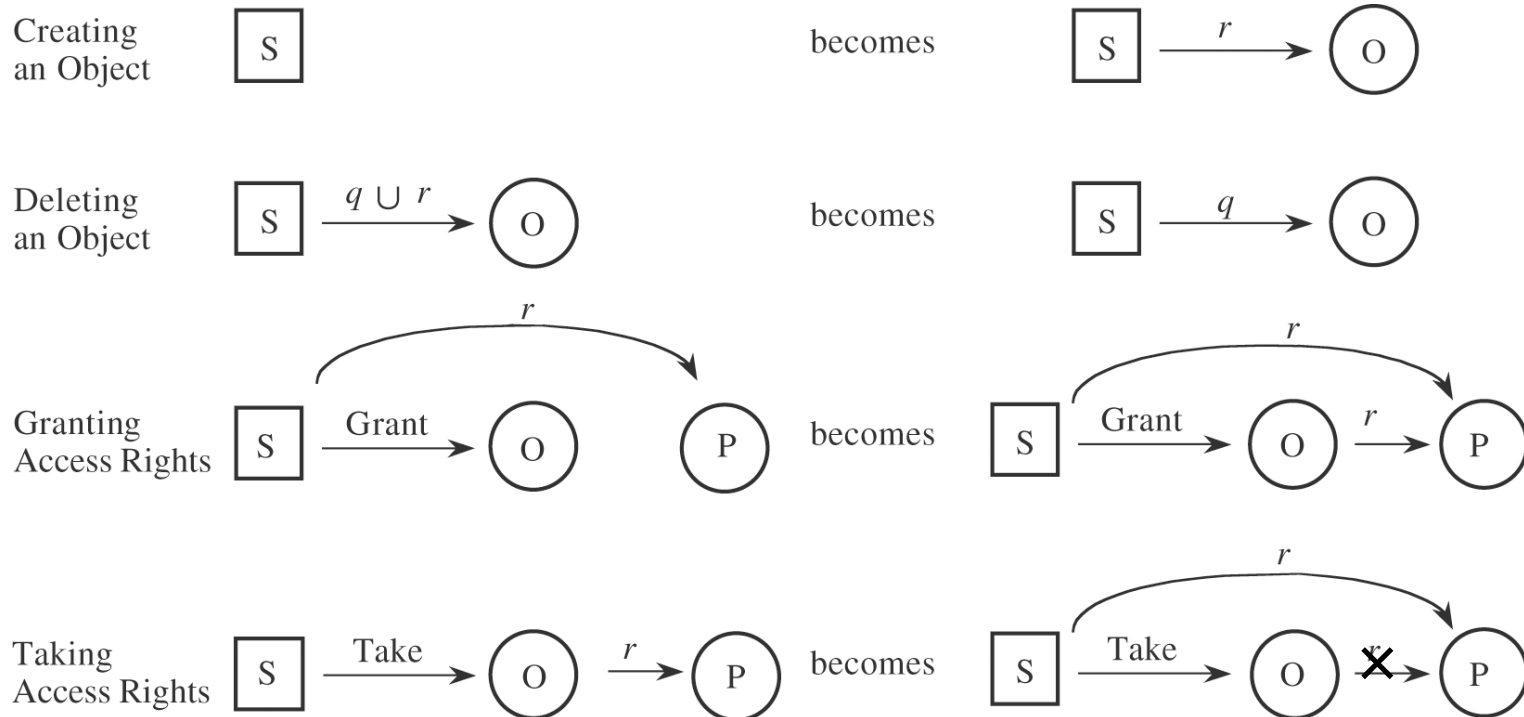


# HRU Results

- In the modeled system, in which commands are restricted to a single operation each, it *is* possible to decide whether a given subject can even obtain a particular right to an object.
- If commands are not restricted to one operation each, it is *not* always decidable whether a given protection system can confer a given right.

# Take-Grant Model

- 4 primitive operations by a subject  $s$ :  $\text{create}(o, r)$ ,  $\text{revoke}(o, r)$ ,  $\text{grant}(o, p, r)$ ,  $\text{take}(o, p, r)$ .



# Take-Grant Results

Decidable on certain protection questions in reasonable time:

1. Can we decide whether a given subject can share an object with another subject?
2. Can we decide whether a given subject can steal access to an object from another subject?

# Design of a Trusted OS

- Principles
- Security features of an ordinary OS
- Security features of a trusted OS
- Kernelized design
- Separation/Isolation
- Virtualization
- (Layered design)

# Design Principles by Saltzer & Schroeder

- Least privilege
- Economy of mechanism (small, simple, straightforward)
- Open design
- Complete mediation
- Permission based (default is denial of access)
- Separation of privilege
- Least common mechanism (minimal sharing)
- Ease of use

# Security Features of a Trusted OS


- User identification and authentication
- Mandatory access control (MAC)
- Discretionary access control (DAC)
- Object reuse protection
- Complete mediation
- Trusted path
- Audit
- Intrusion detection

# Kernelized Design

- A **security kernel** is responsible for enforcing the security mechanisms of the entire OS
  - Coverage (complete mediation)
  - Separation (from the rest of OS and applications)
  - Unity (a single set of code thus easier to debug)
  - Modifiability (easy to modify)
  - Compactness (small)
  - Verifiability (friendly to rigorous analysis)

# Reference Monitor

- The most important part of a security kernel
- It must be:
  - Tamperproof
  - Unbypassable
  - analyzable





Pfleeger/Pfleeger Fig. 05-12 rev 9/30/02



# Trusted Computing Base (TCB)

- Everything in a trusted OS necessary to enforce the security policy
  - Trust in the security of the whole system depends on the TCB
  - Non-TCB part can be malicious w/o affecting security enforcement
- Usually include:
  - Hardware
  - Protected memory
  - Some notion of processes and IPC
  - Primitive files (security-related)

# TCB Implementation



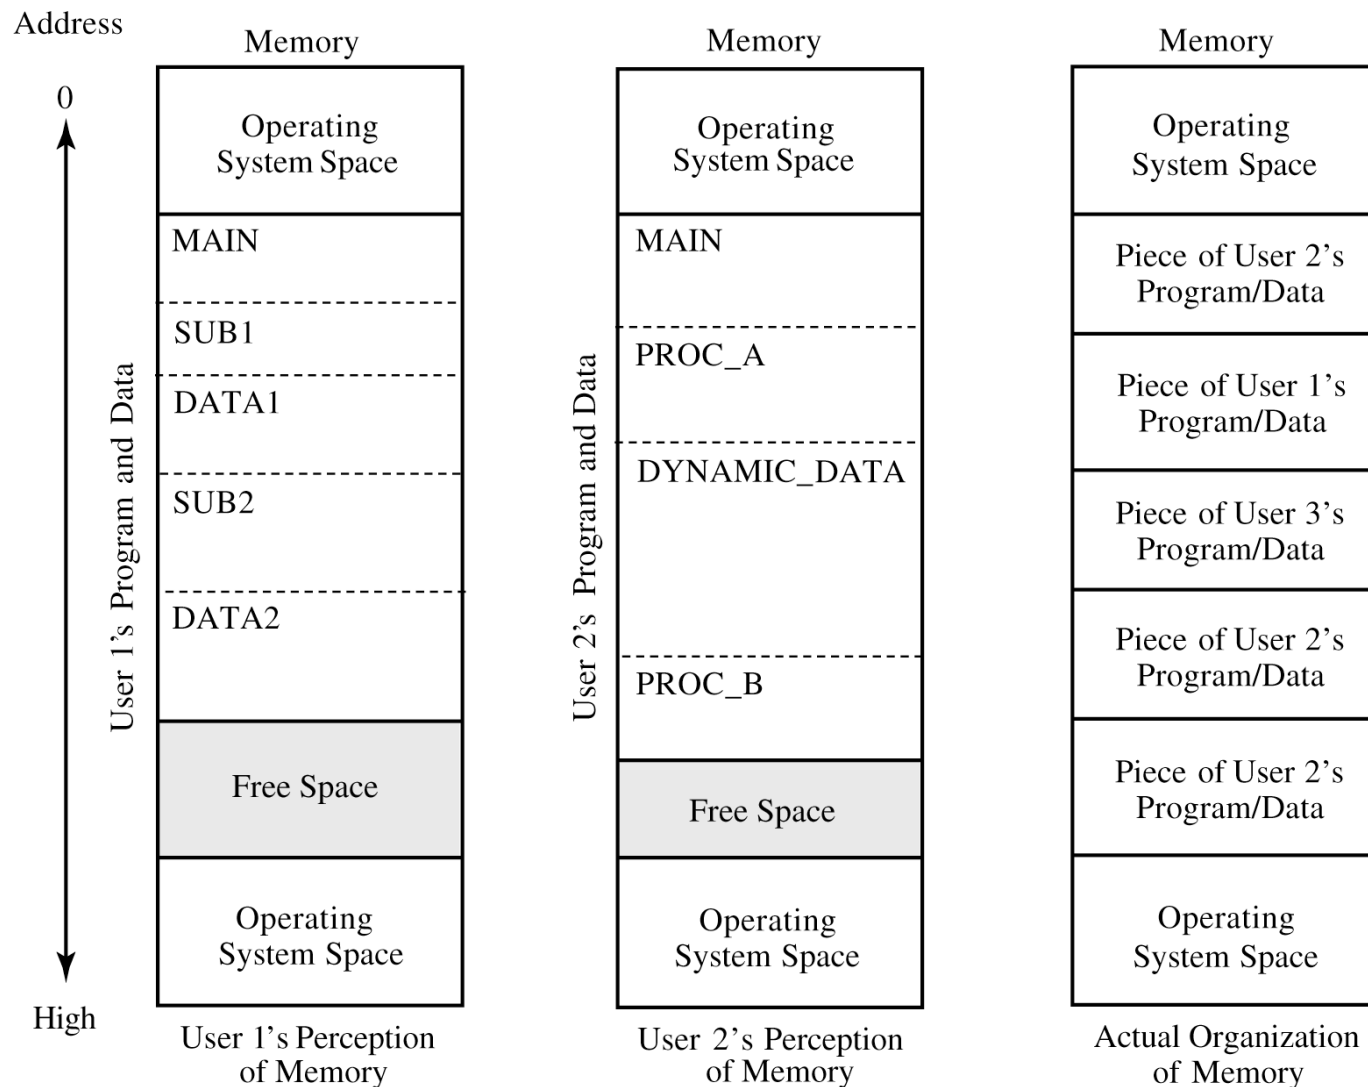
# Separation

- Physical separation
- Temporal separation
- Cryptographic separation
- Logical separation


# Virtualization

- Multiple Virtual Memory Spaces
- Virtual Machines

# Multiple Virtual Memory Spaces



# Virtual Machines



# Virtual Machines

